

```
#!/usr/bin/env python
```

```
"""CEVAE model  
"""
```

```
import edward as ed  
import tensorflow as tf  
import sys
```

```
from edward.models import Bernoulli, Normal  
from progressbar import ETA, Bar, Percentage, ProgressBar
```

```
import numpy as np  
import time  
from scipy.stats import sem  
from sklearn.model_selection import train_test_split  
from argparse import ArgumentParser
```

```
from tensorflow.contrib.layers.python.layers import initializers  
from tensorflow.contrib import slim
```

```
from generate_data import gen_lrnf, ampute, gen_dlvf
```

```
def fc_net(inp, layers, out_layers, scope, lamba=1e-3, activation=tf.nn.relu, reuse=None,  
           weights_initializer=initializers.xavier_initializer(uniform=False)):
```

```
    # utils for cevea
```

```
    with slim.arg_scope([slim.fully_connected,  
                        activation_fn=activation,  
                        normalizer_fn=None,  
                        weights_initializer=weights_initializer,  
                        reuse=reuse,  
                        weights_regularizer=slim.l2_regularizer(lamba)):
```

```
        if layers:
```

```
            h = slim.stack(inp, slim.fully_connected, layers, scope=scope)
```

```
            if not out_layers:
```

```
                return h
```

```
        else:
```

```
            h = inp
```

```
        outputs = []
```

```
        for i, (outdim, activation) in enumerate(out_layers):
```

```
            o1 = slim.fully_connected(h, outdim, activation_fn=activation, scope=scope +
```

```
            '_{}'.format(i + 1))
```

```
            outputs.append(o1)
```

```
        return outputs if len(outputs) > 1 else outputs[0]
```

```
def get_y0_y1(sess, y, f0, f1, shape=(), L=1, verbose=True):
```

```
    # utils for cevea
```

```
    y0, y1 = np.zeros(shape, dtype=np.float32), np.zeros(shape, dtype=np.float32)
```

```
    ymean = y.mean()
```

```
    for l in range(L):
```

```
        if L > 1 and verbose:
```

```
            sys.stdout.write('\r Sample {}/{}'.format(l + 1, L))
```

```
            sys.stdout.flush()
```

```
            y0 += sess.run(ymean, feed_dict=f0) / L
```

```
            y1 += sess.run(ymean, feed_dict=f1) / L
```

```
    if L > 1 and verbose:
```

```
        print()
```

```
    return y0, y1
```

```
def cevae_tf(X, T, Y, n_epochs=100, early_stop = 10, d_cevae=20):
```

```
    T, Y = T.reshape((-1,1)), Y.reshape((-1,1))
```

```
    args = dict()
```

```
    args['earl'] = early_stop
```

```
    args['lr'] = 0.001
```

```
    args['opt'] = 'adam'
```

```

args['epochs'] = n_epochs
args['print_every'] = 10
args['true_post'] = True

M = None # batch size during training
d = d_cevea # latent dimension
lambda = 1e-4 # weight decay
nh, h = 3, 200 # number and size of hidden layers

contfeats = list(range(X.shape[1])) # all continuous
binfeats = []

# need for early stopping
xtr, xva, ttr, tva, ytr, yva = train_test_split(X, T, Y)

# zero mean, unit variance for y during training
ym, ys = np.mean(Y), np.std(Y)
ytr, yva = (ytr - ym) / ys, (yva - ym) / ys
best_logpvalid = - np.inf

with tf.Graph().as_default():
    sess = tf.InteractiveSession()

    ed.set_seed(1)
    np.random.seed(1)
    tf.set_random_seed(1)

    # x_ph_bin = tf.placeholder(tf.float32, [M, len(binfeats)], name='x_bin') # binary inputs
    x_ph_cont = tf.placeholder(tf.float32, [M, len(contfeats)], name='x_cont') # continuous inputs
    t_ph = tf.placeholder(tf.float32, [M, 1])
    y_ph = tf.placeholder(tf.float32, [M, 1])

    # x_ph = tf.concat([x_ph_bin, x_ph_cont], 1)
    x_ph = x_ph_cont
    activation = tf.nn.elu

    # CEVAE model (decoder)
    # p(z)
    z = Normal(loc=tf.zeros([tf.shape(x_ph)[0], d]), scale=tf.ones([tf.shape(x_ph)[0], d]))

    # p(x|z)
    hx = fc_net(z, (nh - 1) * [h], [], 'px_z_shared', lambda=lambda, activation=activation)

    # logits = fc_net(hx, [h], [[len(binfeats), None]], 'px_z_bin', lambda=lambda, activation=activation)
    # x1 = Bernoulli(logits=logits, dtype=tf.float32, name='bernoulli_px_z')

    mu, sigma = fc_net(hx, [h], [[len(contfeats), None], [len(contfeats), tf.nn.softplus]], 'px_z_cont', lambda=lambda,
        activation=activation)
    x2 = Normal(loc=mu, scale=sigma, name='gaussian_px_z')

    # p(t|z)
    logits = fc_net(z, [h], [[1, None]], 'pt_z', lambda=lambda, activation=activation)
    t = Bernoulli(logits=logits, dtype=tf.float32)

    # p(y|t, z)
    mu2_t0 = fc_net(z, nh * [h], [[1, None]], 'py_t0z', lambda=lambda, activation=activation)
    mu2_t1 = fc_net(z, nh * [h], [[1, None]], 'py_t1z', lambda=lambda, activation=activation)
    y = Normal(loc=t * mu2_t1 + (1. - t) * mu2_t0, scale=tf.ones_like(mu2_t0))

    # CEVAE variational approximation (encoder)
    # q(t|x)
    logits_t = fc_net(x_ph, [d], [[1, None]], 'qt', lambda=lambda, activation=activation)

```

```

qt = Bernoulli(logits=logits_t, dtype=tf.float32)
# q(y|x,t)
hqy = fc_net(x_ph, (nh - 1) * [h], [], 'qy_xt_shared', lambda=lambda, activation=activation)
mu_qy_t0 = fc_net(hqy, [h], [[1, None]], 'qy_xt0', lambda=lambda, activation=activation)
mu_qy_t1 = fc_net(hqy, [h], [[1, None]], 'qy_xt1', lambda=lambda, activation=activation)
qy = Normal(loc=qt * mu_qy_t1 + (1. - qt) * mu_qy_t0, scale=tf.ones_like(mu_qy_t0))

# q(z|x,t,y)
inpt2 = tf.concat([x_ph, qy], 1)
hqz = fc_net(inpt2, (nh - 1) * [h], [], 'qz_xty_shared', lambda=lambda, activation=activation)
muq_t0, sigmaq_t0 = fc_net(hqz, [h], [[d, None], [d, tf.nn.softplus]], 'qz_xt0', lambda=lambda,
activation=activation)
muq_t1, sigmaq_t1 = fc_net(hqz, [h], [[d, None], [d, tf.nn.softplus]], 'qz_xt1', lambda=lambda,
activation=activation)
qz = Normal(loc=qt * muq_t1 + (1. - qt) * muq_t0, scale=qt * sigmaq_t1 + (1. - qt) * sigmaq_t0)

# Create data dictionary for edward
data = {x2: x_ph_cont, y: y_ph, qt: t_ph, t: t_ph, qy: y_ph}

# sample posterior predictive for p(y|z,t)
y_post = ed.copy(y, {z: qz, t: t_ph}, scope='y_post')
# crude approximation of the above
y_post_mean = ed.copy(y, {z: qz.mean(), t: t_ph}, scope='y_post_mean')
# construct a deterministic version (i.e. use the mean of the approximate posterior) of the lower bound
# for early stopping according to a validation set
y_post_eval = ed.copy(y, {z: qz.mean(), qt: t_ph, qy: y_ph, t: t_ph}, scope='y_post_eval')
# x1_post_eval = ed.copy(x1, {z: qz.mean(), qt: t_ph, qy: y_ph}, scope='x1_post_eval')
x2_post_eval = ed.copy(x2, {z: qz.mean(), qt: t_ph, qy: y_ph}, scope='x2_post_eval')
t_post_eval = ed.copy(t, {z: qz.mean(), qt: t_ph, qy: y_ph}, scope='t_post_eval')
logp_valid = tf.reduce_mean(tf.reduce_sum(y_post_eval.log_prob(y_ph) + t_post_eval.log_prob(t_ph), axis=1) +
tf.reduce_sum(x2_post_eval.log_prob(x_ph_cont), axis=1) +
tf.reduce_sum(z.log_prob(qz.mean()) - qz.log_prob(qz.mean()), axis=1))

inference = ed.KLqp({z: qz}, data)
optimizer = tf.train.AdamOptimizer(learning_rate=args['lr'])
inference.initialize(optimizer=optimizer)

saver = tf.train.Saver(tf.contrib.slim.get_variables())
tf.global_variables_initializer().run()

n_epoch, n_iter_per_epoch, idx = args['epochs'], 10 * int(xtr.shape[0] / 100), np.arange(xtr.shape[0])

# # dictionaries needed for evaluation
t0, t1 = np.zeros((X.shape[0], 1)), np.ones((X.shape[0], 1))
# tr0t, tr1t = np.zeros((xte.shape[0], 1)), np.ones((xte.shape[0], 1))
f1 = {x_ph_cont: X, t_ph: t1}
f0 = {x_ph_cont: X, t_ph: t0}
# f1t = {x_ph_bin: xte[:, 0:len(binfeats)], x_ph_cont: xte[:, len(binfeats):], t_ph: tr1t}
# f0t = {x_ph_bin: xte[:, 0:len(binfeats)], x_ph_cont: xte[:, len(binfeats):], t_ph: tr0t}

for epoch in range(n_epoch):
    avg_loss = 0.0

```

```
widgets = ["epoch #%d|" % epoch, Percentage(), Bar(), ETA()]
pbar = ProgressBar(n_iter_per_epoch, widgets=widgets)
pbar.start()
np.random.shuffle(idx)
for j in range(n_iter_per_epoch):
    # print('j', j)
    # pbar.update(j)
    batch = np.random.choice(idx, 100)
    x_train, y_train, t_train = xtr[batch], ytr[batch], ttr[batch]
    info_dict = inference.update(feed_dict={x_ph_cont: x_train,
                                             t_ph: t_train, y_ph: y_train})

    avg_loss += info_dict['loss']

avg_loss = avg_loss / n_iter_per_epoch
avg_loss = avg_loss / 100

if epoch % args['earl'] == 0 or epoch == (n_epoch - 1):
    logpvalid = sess.run(logp_valid, feed_dict={x_ph_cont: xva,
                                                t_ph: tva, y_ph: yva})

    if logpvalid >= best_logpvalid:
        print('Improved validation bound, old: {:.3f}, new: {:.3f}'.format(
best_logpvalid, logpvalid))
        best_logpvalid = logpvalid
        saver.save(sess, 'data/cevea_models/dlvm')

    saver.restore(sess, 'data/cevea_models/dlvm')
    y0, y1 = get_y0_y1(sess, y_post, f0, f1, shape=Y.shape, L=100)
    y0, y1 = y0 * ys + ym, y1 * ys + ym

sess.close()

return y0.reshape((-1)), y1.reshape((-1))

if __name__ == '__main__':
    Z, X, w, y, ps = gen_dlv()
    y0, y1 = cevae_tf(X, w, y, n_epochs=10)
    print('cevae_tf OKAY !')
    print(y0.shape, y1.shape)
```